

Lecture 4:

A* Search, Minimax search and Decision Trees

30 January 2020

How does GD algorithm know it has reached min/max?


- $\nabla h = 0$
 - Could be a min, max or saddle point

How to determine if min/max

- 1d - second derivative test
 - $\nabla^2 h > 0 \Rightarrow$ local min
 - $\nabla^2 h < 0 \Rightarrow$ local max
 - $\nabla^2 h = 0 \Rightarrow$ undetermined, may be saddle point
- Higher dimensions, use Hessian to perform second derivative test

$$\bullet \quad \mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$

- **No, you don't need to know this!!!**



- Order nodes by sum of path length and heuristic
- Priority queue order by $\text{path}(s) + h(s)$
- Will find optimal path if using admissible heuristic

```
priority_queue.push(initial state, 0)
```

```
While(!priority_queue.empty())
```

```
    s = priority_queue.pop()
```

```
    if(s == goal)
```

```
        return
```

```
    if (!s.visited)
```


- $s.\text{visited} = \text{true}$

- for all actions $a \in A$

```
     $s' = \tau(s, a)$ 
```

```
     $s'_{\text{path\_length}} = s_{\text{path\_length}} + \text{length}(a)$ 
```

```
    priority_queue.push( $s'$ ,  $s'_{\text{path\_length}} + h(s')$ )
```




- A heuristic is admissible if it never overestimates the cost from a node to the goal
 - $h(s) \leq \text{cost}(\text{path}(s, s_{\text{goal}}))$
 - Multiple goals or paths to goal
 - $h(s)$ less than minimum cost over all paths to goals
 - Example: Find a path to the nearest post office
 - Doesn't matter what post office you get to
 - $h(s)$ must be less than cost of shortest path to nearest post office
- Examples
 - Robot in plane
 - Euclidean distance
 - Robot in grid
 - Manhattan distance
 - Mahjong Solitaire
 - $1/2 * \text{number of tiles left}$



- Proof that search with admissible heuristic will always return optimal path to goal
- Assumptions
 - All actions have positive cost
 - $h(s)$ is admissible
 - $h(s) \leq \text{cost}(\text{path}(s, s_{\text{goal}}))$
 - Search expands nodes using priority queue ordered by $g(s)$
 - $g(s) = \text{cost}(\text{path}(s_{\text{initial}}, s)) + h(s)$

- Map of driving routes
 - Find shortest route
 - Heuristic = Direct Euclidean distance to goal



City	$h(s)$ = direct distance to Amarillo
Wichita Falls	220
Dallas	330
Austin	505
Houston	605
San Antonio	445
San Angelo	315
El Paso	440
Laredo	670
Brownsville	795

How can we apply A* to problems with continuous state spaces?

- Robot in XY-plane

- Generate graph in space of problem
- Grid up state space
 - What resolution to use
 - More cells \Rightarrow higher computation cost
 - Complexity exponential with respect to dimension of state space
 - $O(x^d)$
 - Curse of dimensionality
- Generate graph in space of problem
 - Example robotics
- Adaptive grid
 - Recursively subdivide cells that are interesting
 - Example: Robotics - use higher resolution cells in difficult regions

