


Lecture 7:


Decision Trees and Inductive Learning

11 February 2020



- Look ahead d moves
- Compute h of all leaves
- Inductively compute values for internal nodes from values of children
 - Max Step (agent's turn): Set values to be max value of children
 - Min Step (opponent's turn): Set value to min value of children
- Select child of root that has highest value
 - This is the action the agent should make

MIN



- Expand nodes in depth first ordering
- Expand action c_1 and get 2
- What can you say about about state C and action a_2 ?
 - At state C opponent will be able to get to a state with a value ≤ 2
 - Already know that a_1 gives you a value of 3
 - Never want to take action a_2
- Don't even need to consider other children of C

- Nodes are states
 - At each state you need to make a decision specific to that state
 - Edge for each choice
 - Edge leads to state obtained by selecting choice
- Effectively, actions are choices
 - Different from other examples because actions are different for each state
- Example - navigating a road map
 - At intersection of Main and Oak
 - Turn right on Oak
 - Turn Left on Oak
 - Stay on Main Street
 - Options available different for different intersections



- Problem
 - Given a list of choices construct a decision tree
 - Given a list of restaurants
 - Construct a decision tree
 - Decide what type of restaurant to go to



	Type	Price	Crowded	Time
McDonald's	Fast Food	\$	Yes	< 10 min
Burger King	Fast Food	\$	Yes	< 10 min
Chipotle	Mexican	\$\$	Yes	< 10 min
Jennie's	Diner	\$\$\$	Yes	< 10 min
Fred's Fried Fish	Seafood	\$	No	< 10 min
Captain John's	Seafood	\$\$\$	No	20 min
Henri's	French	\$\$\$	No	30 min
Pinocchio's	Italian	\$\$	Yes	20 min
Taste of Italy	Italian	\$\$	No	20 min


- We want to grow a simple tree
 - a good attribute prefers attributes that split the data so that each successor node is as *pure* as possible
 - i.e., the distribution of examples in each node is so that it mostly contains examples of a single class
- In other words:
 - We want a measure that prefers attributes that have a high degree of „order“:
 - Maximum order: All examples are of the same class
 - Minimum order: All classes are equally likely
 - **Entropy** is a measure for (un-)orderedness
 - Another interpretation:
 - Entropy is the amount of information that is contained
 - all examples of the same class → no information

- S is a set of examples
- p_{\oplus} is the proportion of examples in class \oplus
- $p_{\ominus} = 1 - p_{\oplus}$ is the proportion of examples in class \ominus

Entropy:

$$E(S) = -p_{\oplus} \cdot \log_2 p_{\oplus} - p_{\ominus} \cdot \log_2 p_{\ominus}$$

- Interpretation:
 - amount of unorderedness in the class distribution of S



- For a complex (non-binary) set Entropy is given by the following

$$E(S) = \sum_{x \in X} -p(x) \log_2 p(x)$$

- **Problem:**

- Entropy only computes the quality of a single (sub-)set of examples
 - corresponds to a single value
- How can we compute the quality of the entire split?
 - corresponds to an entire attribute

- **Solution:**

- Compute the weighted average over all sets resulting from the split
 - weighted by their size

$$I(S, A) = \sum_i \frac{|S_i|}{|S|} \cdot E(S_i)$$

- **Example:**

- Average entropy for attribute *Outlook*:

$$I(\text{Outlook}) = \frac{5}{14} \cdot 0.971 + \frac{4}{14} \cdot 0 + \frac{5}{14} \cdot 0.971 = 0.693$$

- When an attribute A splits the set S into subsets S_i
 - we compute the average entropy
 - and compare the sum to the entropy of the original set S

Information Gain for Attribute A

$$\text{Gain}(S, A) = E(S) - I(S, A) = E(S) - \sum_i \frac{|S_i|}{|S|} \cdot E(S_i)$$

- The attribute that maximizes the difference is selected
 - i.e., the attribute that reduces the unorderedness most!
- **Note:**
 - maximizing information gain is equivalent to minimizing average entropy, because $E(S)$ is constant for all attributes A

- Greedy approach
 - Recursively select trait that maximizes information gain.

Build_Decision_Tree(S , Attributes)

- If S is monotone (all same type) or Attributes = {}
 - return
- //Find best attribute
- $IG_{best} = -1$
- for all attributes a
 - $S' =$ set of subsets from dividing S by a
 - if $IG(S') > IG_{best}$
 - $IG_{best} = IG(S')$
 - $a_{best} = a$
- $S' =$ set of subsets from dividing S by a_{best}
- For all S' in
 - Build_Decision_Tree(S' , Attributes \ a_{best})

- Recursively subdivide using attribute that maximizes information gain

$$E(S) - \sum_i \frac{|S_i|}{|S|} \cdot E(S_i) \quad \text{where} \quad E(S) = \sum_{x \in X} -p(x) \log_2 p(x)$$



- Intuitively, which division looks the best?

	Type	Price	Crowded	Time
McDonald's	Fast Food	\$	Yes	< 10 min
Burger King	Fast Food	\$	Yes	< 10 min
Chipotle	Mexican	\$\$	Yes	< 10 min
Jennie's	Diner	\$\$\$	Yes	< 10 min
Fred's Fried Fish	Seafood	\$	No	< 10 min
Captain John's	Seafood	\$\$\$	No	20 min
Henri's	French	\$\$\$	No	30 min
Pinocchio's	Italian	\$\$	Yes	20 min
Taste of Italy	Italian	\$\$	No	20 min

$$E(S) - \sum_i \frac{|S_i|}{|S|} \cdot E(S_i)$$

where $E(S) = \sum_{x \in X} -p(x) \log_2 p(x)$

- After divide by price



	Type	Crowded	Time
McDonald's	Fast Food	Yes	< 10 min
Burger King	Fast Food	Yes	< 10 min
Fred's Fried Fish	Seafood	No	< 10 min
Chipotle	Mexican	Yes	< 10 min
Pinocchio's	Italian	Yes	20 min
Taste of Italy	Italian	No	20 min
Jennie's	Diner	Yes	< 10 min
Captain John's	Seafood	No	20 min
Henri's	French	No	30 min