

CS 440

Introduction to Artificial Intelligence

Lecture 22:

Cross-Validation – Linear Regression (continued)


May 9, 2020

- **What is Learning?**
 - Use previous experience to solve problem
 - Example solutions to problems
 - Learning by demonstration
 - Examine trends in data to predict solution
 - Data mining
 - Train agent to perform task
 - Reinforcement learning

- **Model**
 - Representation of environment
 - Queried to find solutions to problems
- **Models used in first two thirds of class**
 - **Examples:**
 - State/Action/Transition/Reward models
 - Bayesian networks
 - Markov models
 - Logical statements
 - **We defined these models**
 - **Agent used models we built**
- **What if we allowed the agent to build the models?**

- **Allow agent to build or modify model**
 - **Example: Robot map its environment**
 - **Robot must generate some representation of its environment**
 - **Able to query this representation**
- **Models may not be intuitive to programmer**
 - **Mapping of inputs to solutions**

A Regression Problem




$$y = f(x) + \text{noise}$$

Can we learn f from this data?

Let's consider three methods...

Linear Regression




$$y = w_0 + w_1 \cdot x$$

Objective: Minimize the
Sum of Squared Errors
i.e. sum of squared differences
between y values
and the green line

- \hat{y}_i is the prediction of the linear model
- y_i the actual value for input x_i
- Then minimize: $Q = \sum_i (\hat{y}_i - y_i)^2$


Quadratic Regression



Objective: Minimize the
Sum of Squared Errors
i.e. sum of squared differences
between y values
and the green curve




$$y = w_0 + w_1 \cdot x + w_2 \cdot x^2$$

Join-the-dots



Also known as **piecewise linear nonparametric regression** if that makes you feel better


y = complicated



Why not choose the method with the best fit to the data?

“How well are you going to predict future data drawn from the same distribution?”


- **How can we evaluate the fidelity of a model?**
 - **Minimize error function**
 - **Lead to over-fitting**
 - **More complicated model**
 - **More expensive to train and query**
 - **Performance levels off and in some cases declines**



(Linear regression example)

Mean Squared Error = 2.4


1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the **test set**



(Quadratic regression example)

Mean Squared Error = 0.9

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the test set



(Join the dots example)

Mean Squared Error = 2.2

1. Randomly choose 30% of the data to be in a **test set**
2. The remainder is a **training set**
3. Perform your regression on the training set
4. Estimate your future performance with the test set

Good news:

- Very very simple
- Can then simply choose the method with the best test-set score


Bad news:


- Wastes data: we get an estimate of the best method to apply to 30% less data
- If we don't have much data, our test-set might just be lucky or unlucky

We say the “test-set estimator of performance has high variance”

For $k=1$ to R


1. Let (x_k, y_k) be the k^{th} record





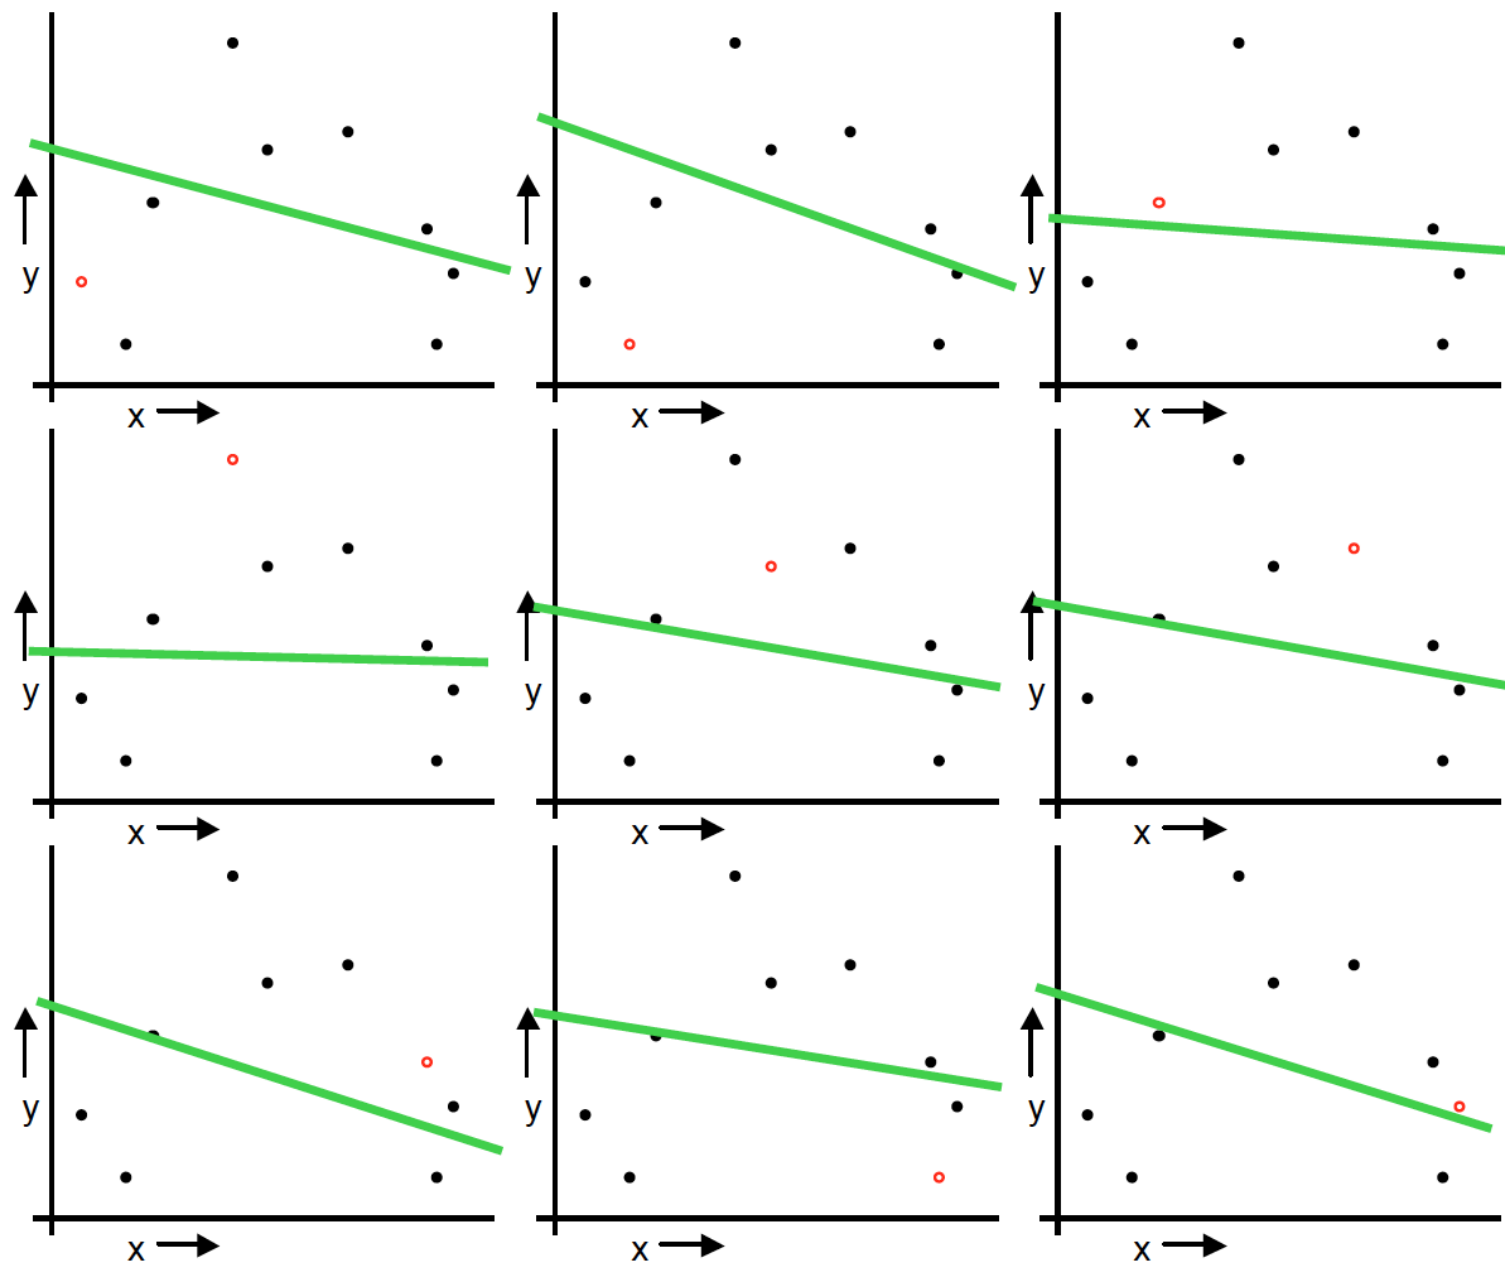
For $k=1$ to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining $R-1$ datapoints



For $k=1$ to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining $R-1$ datapoints
4. Note your error (x_k, y_k)



For $k=1$ to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining $R-1$ datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

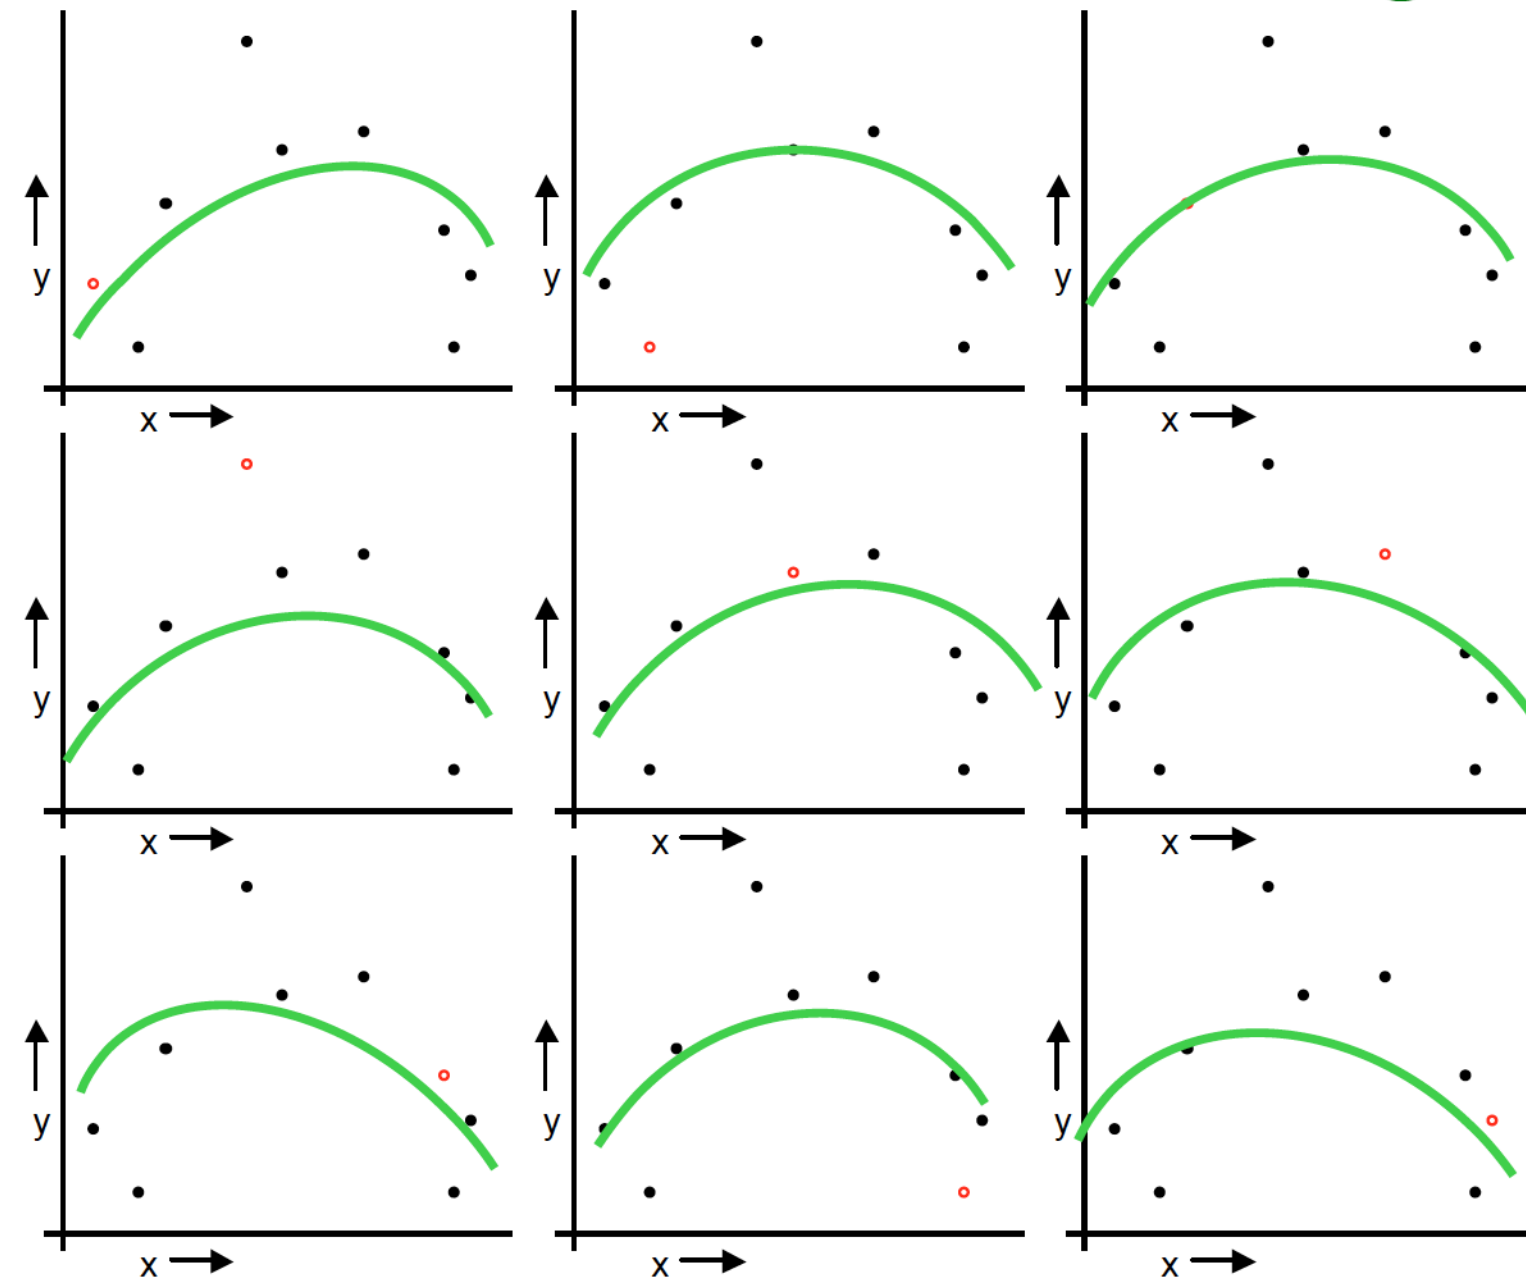
$$MSE_{LOOCV} = 2.12$$


For $k=1$ to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining $R-1$ datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

$$MSE_{LOOCV} = 0.962$$





For $k=1$ to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining $R-1$ datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

$$MSE_{LOOCV} = 3.33$$

	Downside	Upside
Test-set	Variance: unreliable estimate of future performance	Cheap
Leave-one-out	Expensive. Has some weird behavior	Doesn't waste data

k-fold Cross Validation


Randomly break the dataset into k partitions (in our example we'll have k=3 partitions colored Red Green and Blue)

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error



Linear Regression

$$MSE_{3FOLD} = 2.05$$


k-fold Cross Validation

Randomly break the dataset into k partitions (in our example we'll have k=3 partitions colored Red Green and Blue)

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.



Quadratic Regression

$$MSE_{3FOLD} = 1.11$$

Then report the mean error


k-fold Cross Validation

Randomly break the dataset into k partitions (in our example we'll have k=3 partitions colored Red Green and Blue)

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.



Joint-the-dots
 $MSE_{3FOLD} = 2.93$

Then report the mean error